[00:00:20.760]
Welcome to the Making Sense podcast, this is Sam Harris. OK, so today I'm bringing you a conversation that I originally recorded for the Waking Up app and we released it, there's a series of separate lessons a couple of weeks back, but the response has been such that I wanted to share it here on the podcast and put it outside the paywall. This seems like a better holiday message than most. As I think many of you know, waking up isn't just a meditation app at this point, it's really the place where I do most of my thinking about what it means to live a good life.

[00:01:00.180]
And this conversation is about generosity and about how we should think about doing good in the world. Increasingly, I'm looking to use this podcast and the waking up app to do more than merely spread what I consider to be good ideas, that's their primary purpose, obviously. But I want to help solve some of the worst problems we face more directly than just talking about them. And I want to do this systematically, really thinking through what it takes to save the most lives or reduce the worst suffering or mitigate the most catastrophic risks.

[00:01:38.580]
And to this end, I've taken the pledge over it, given what we can, which is the foundation on effective altruism started by the philosophers Will McCaskill and Toby Ord, both of whom have been on the podcast. And this pledge to give a minimum of 10 percent of one's pre-tax income to the most effective charities. I've also taken the founder's pledge, which amounts to the same thing, and I've had waking up become one of the first corporations to pledge a minimum of 10 percent of its profits to charity.

[00:02:09.800]
And the thinking behind all of this is the subject of today's podcast. Of course, there is a bias against speaking about this sort of thing in public or even in private, right. It's often believed that it's better to practice one's generosity anonymously because then you can be sure you're doing it for the right reasons. You're not trying to just burnish your reputation. As you'll hear in today's conversation, there are very good reasons to believe that this is just not true and that the imagined moral virtue of anonymity is something we really need to rethink.

[00:02:45.340]
In fact, I've just learned of the knock on effects of the few times I have discussed my giving to charity on this podcast and their surprisingly substantial. Just to give you a sense of it, last year I released an episode titled Knowledge and Redemption, where we discuss the Bard Prison Initiative based on the PBS documentary that Lynn Novick and Ken Burns did. And Lynn was on that podcast. And at the end, I think I asked you all to consider supporting that work, too.

[00:03:17.650]
And together we donated one hundred and fifty thousand dollars based on that one episode alone. I've also occasionally mentioned on the podcast that I donate each month to the Against Malaria Foundation and was actually my first podcast conversation with Wil MacAskill that convinced me to do that. I do through the charity evaluator give well, Doug well, the good people to give well, just told me that they've received over five hundred thousand dollars in donations from you guys and they expect another five hundred thousand dollars over the next year from podcast listeners who have set up their donations on a recurring basis.

[00:03:55.690]
So that's one million dollars and many lives saved just as a result of some passing comments I've made on the podcast. And then I've heard from Wil McCaskill's people over it, given what we can wear, I took their 10 percent pledge, which I haven't spoken about much, but it seems that hundreds of you have also taken that pledge, again, unsolicited by me, but specifically attributing this podcast and the waking up app as the reason that's hundreds of people, some of whom may be quite wealthy or will become wealthy, who have now publicly pledged to give a minimum of 10 percent of their pre-tax income to the most effective charities every year for the rest of their lives.

[00:04:47.450]
That is awesome. So all of this inspired me to share this conversation from the waking up app, again, this is a fairly structured conversation with the philosopher Wil MacAskill. Some of you may remember the conversation I had with Will four years ago on the podcast that was episode number 44. And that's a great companion to today's episode because it gets into some of the fundamental issues of ethics here. Today's conversation is much more focused on the actions we can all take to make the world better.

[00:05:20.260]
And how we should think about doing that Will and I challenge some old ideas around giving and we discuss why they're really not very good ideas in the end. You'll also hear that there's still a lot of moral philosophy to be done in this area. I don't think these issues are fully worked out at all, and that's really exciting. There's a lot to talk about here and there's something for moral philosophers to actually do that might really matter to the future of our species.

[00:05:49.440]
In particular, I think there's a lot of work to be done on the ethics of wealth inequality, both globally and within the wealthiest societies themselves, and I'm sure I will do many more podcasts on this topic. I suspect that wealth inequality is producing much, if not most, of our political conflict at this point. And it certainly determines what we do with our resources. So I think it's one of the most important topics of our time anyway, Will and I cover a lot here, including how to choose causes to support and how best to think about choosing a career so as to do the most good over the course of one's life.

[00:06:27.120]
The question that underlies all of this really is how can we live a morally beautiful life, which is more and more what I care about, in which the young will McCaskill is certainly doing, as you will hear. Finally, I want to again recognize all of you who have made these donations and pledges, as well as the many of you who have been supporting my work these many years, and also the many of you who have become subscribers to the podcast in the last year.

[00:06:58.020]
I couldn't be doing any of these things without you. And I certainly look forward to what we're going to do next. Twenty twenty one should be an interesting here. So my deep thanks to all of you. And now I bring you well, MacAskill. I am here with Wil MacAskill. Well, thanks for joining me again. Thanks so much for having me on. So I just posted a conversation and you and I had four years ago on my podcast, Anta waking up as well, because I thought it was such a useful introduction to many of the issues we're going to talk about.

[00:07:39.230]
And it's a different conversation because we got into very interesting questions of moral philosophy that I think we probably won't focus on here. So it just seems like a great background for the series of lessons we're now going to sketch out in a conversation. But for those who have not taken the time to listen to that just yet, maybe we should summarize your background here. Who are you? Well, and how do you come to have any opinion about altruism, generosity and what it means to live a good life?

[00:08:10.310]
Give us your potted bio.

[00:08:12.230]
So, yeah, my potted bio. So I grew up in Glasgow and I was always interested in two things. One was kind of ideas and then in particular philosophy when I discovered that and second was interested in helping people. So as a teenager, I volunteered running summer camps for children with who impoverished and had disabilities.

[00:08:34.490]

I worked at kind of old folks home, but then it was when I came across the arguments of Peter Singer in particular, his arguments that we have the moral obligation to be giving away most of our income to help people in very poor countries simply because such a move would not be a great burden on us.

[00:08:53.090]
It would be a financial sacrifice, but not an enormous sacrifice in terms of our quality of life, but could make an enormous difference for hundreds of people around the world. That moved me very much. But kind of being human, I didn't really do very much on the basis of those arguments for many years until I came to Oxford, met another philosophical, Toby Ord, who has actually very similar ideas and was planning to give away most of his income over the course of his life.

[00:09:20.270]
And together we set up an organisation called Giving What We Can, which encourage people to give at least 10 percent of their income to those organisations. They think that can do the most good. Sam, I know that you have now taken that 10 percent pledge, and I'm delighted that that's the case. And since then, this kind of set of ideas that we're really just to, you know, very impractical philosophy, grad students kind of setting this up and not saying, you know, I certainly never thought it was going to be that big a deal.

[00:09:48.650]
I was just doing it because I thought it was morally very important. It turned out just a lot of people had had similar sets of ideas. And given what we can act like a bit of a lightning about for people all around the world who were motivated to try to do good, but also to do it as effectively as possible, because at the time we had a set of recommended charities, there was also the organisation Give Well whose work. We leaned extremely heavily on making recommendations about what charities that they thought would do the most good and effective altruism at the time, focused on charity in particular, and in particular focused on doing good for people in extreme poverty.

[00:10:27.500]
And since then it's broadened out a lot. So now most people in the effective altruism community, when they're trying to do good, are doing so via their career in particular.

[00:10:37.550]
And there's a much broader range, of course, areas. So animal welfare is a big focus. And in particular, and I think increasing are issues that might potentially affect future generations in a really big way and in particular kind of risks to the future of civilisation at all. Toby, he talks about when he was on your podcast. And I have a factoid in my memory which I think I got from your original interview with Tim Ferriss on his podcast, am I correct in thinking that you were the youngest philosophy professor at Oxford?

[00:11:12.000]
Yes, the precise fact is when I joined the faculty at Oxford, which is age 28, I'm pretty confident I was the youngest associate professor of philosophy in the world at the time.

[00:11:22.590]
Oh, nice. That's all right. No doubt your quickly aging out of that distinction, have you have you lost your record yet?

[00:11:30.100]
Yeah, well, I'm old man at 33 years old now, and I definitely lost that a few years ago.

[00:11:35.260]
Well, so it's great to talk to you about these things because, you know, as you know, you've been very influential in my thinking. You directly inspired me to start giving a minimum of 10 percent of my income to charity and also to commit waking up as a company to give a minimum of 10 percent of its profits to charity. But I'm very eager to have this conversation because it still seems to me there's a lot of thinking yet to do about how to approach doing good in the world.

[00:12:03.860]
There may be some principles that you and I either disagree about or maybe we all agree that we just don't have good enough intuitions to have a strong opinion one way or another. But it really just seems to me to be territory that can benefit from new ideas and new intuition pumps. And it's just a lot to be sorted out here. And I think, you know, as I said, we will have a structured conversation here which will break into a series of lessons.

[00:12:31.600]
And so this is really an introduction to the conversation that's coming. And all of this relates specifically to this movement. You started effective altruism, and we'll get very clear about what that means and what it may yet mean. But this does connect to deeper and broader questions like how should we think about doing good in the world in general and what would it mean to do as much good as possible? And how do those questions connect to questions like what sort of person should I be or what does it mean to live a truly good life?

[00:13:12.250]
These are questions that lie at the core of moral philosophy and at the core of any person's individual attempt to live and examined life and develop an ethical code and just form a vision of what would be a good society. I mean, we're all personally attempting to improve our lives, but we're also trying to converge on a common picture of what it would mean for us to be building a world that is making it more and more likely that humanity is moving in the right direction.

[00:13:48.560]
We have to have a concept of what the goal is here or what what a range of suitable goals might be. And we have to have a concept of when we're wandering into moral error, you know, personally and collectively. So there's a lot to talk about here. And talking about the specific act of trying to help people trying to do good in the world really sharpens up our sense of the stakes here and the opportunities. So I'm really happy to be getting into this with you.

[00:14:19.490]
Before we get into what effective altruism is, I think we should address a basic skepticism that people have and even very rich people have, perhaps especially rich people have. This is a skepticism about altruism itself and in particular, skepticism about charity. And I think there are some good reasons to be skeptical about charity, at least in a in a local context. And then there's some very bad reasons. And I just want to lob you some of these these reasons and we can talk about them because, I mean, I would imagine you've encountered this yourself.

[00:15:00.500] - Sam Harris
I meet some very fortunate people who have immense resources and can do a lot of good in the world, who are fundamentally skeptical about giving to charity. And the bad reason here that I always encounter is something we might call the myth of the self-made man. The idea that there's some it's somehow an ethically impregnable position to notice all the ways in which you are responsible for all of your good luck. No matter how distorted this this appraisal might be. You weren't born into wealth and you made it all yourself and you don't owe anyone anything.

[00:15:42.140] - Sam Harris
And in fact, giving people less fortunate than yourself, any of the resources you've acquired is not really helping them in the end. I mean, you want to teach people to fish, but you don't want to give them fish. There's some Ayn Rand an ethic of radical selfishness combined with a vision of capitalism that, you know, wherein free markets can account for, you know, every human problem simply by all of us behaving like atomized selves, seeking our own our own happiness.

[00:16:14.180] - Sam Harris
We will be no surprise to people who listen to me that I think there's something deeply flawed in this analysis. But what do you do when someone hits you with this ethical argument that they're self-made and everyone should aspire to also pull themselves up by their own bootstraps? And we falsify

something about the project of living a good life by even thinking in terms of altruism and charity.

[00:16:42.360] - Will MacAskill
I think there's a few things to say here. So in the first case, the fact that you're a self-made man, I mean, I do disagree with the premise.

[00:16:51.930] - Will MacAskill
I can predict 80 percent of the of the information about your income just from your place of birth, whereas, you know, you could be the hardest working Bangladeshi in the world. But if you're born into extreme poverty in Bangladesh, it's going to be very difficult indeed to become a billionaire. So I agree with you that that's a myth. But even if we accepted that, the fact that you have rightly earned your money yourself doesn't mean that you don't have any obligations to help other people.

[00:17:20.160] - Will MacAskill
So Peter Singer's now very famous thought experiment. You walk past a pond, it's a very shallow pond, you could easily kind of wade in as deep as you, as deep as you like. And you can see that there's a child standing there now, perhaps it's the case that you have an entirely self-made man, perhaps it's the case that the suit that you wore, you just bought yourself. But that seems neither here nor there with respect to whether you ought to try and wade in and save this child who might be drowning.

[00:17:52.160] - Will MacAskill
And I think that's just quite an intuitive position. In fact, this ideal of self-actualization of kind of being the best version of yourself that you can be, which is the kind of admirable version of this otherwise sometimes quite dark perspective on the world. I think there is like part of being a self-actualized, authentically living person is living up to your ideals and principles.

[00:18:15.920] - Will MacAskill
And for most people in the world, you actually want to be helpful. Altruistic person acting in that way is acting in accordance with your deepest values. That is acting an authentic inner self-actualized life.

[00:18:27.680] - Will MacAskill
And then just on the second point is about whether, well, maybe charity gets in the way. Maybe it's actually harmful because it makes people rely on bailouts. Well, here we've got to just think about, you know, there is market failure where in the case of public goods or externalities, markets don't do what they ought to do.

[00:18:50.150] - Will MacAskill
And perhaps you want government to step in and provide police or defence and street lights or taxes against climate change. And even the most kind of hardcore libertarian free market proponent should accept that's a good thing to do sometimes. But then there's also cases of Democratic failure, too. So what if the potential people are not protected by a functioning democratic governments? That's true for people in poor countries. That's true for non-human animals. That's true for people who are yet to be born, people who don't have a vote – the future generations are disenfranchised.

[00:19:24.590] - Will MacAskill
So we shouldn't expect markets or government to be taking appropriate care of those individuals who are disenfranchised by both the market and by even democratic institutions. And so what else is there apart from philanthropy?

[00:19:38.500] - Sam Harris
Yeah. Yeah. So I've spoken to a lot about the myth of the self-made man. Whenever I criticize, the notion of free will is just obvious that however self-made you are, you didn't create the tools by which you made yourself right. So if you are incredibly intelligent or have an immense capacity for effort, you didn't create any of that about yourself.

[00:20:03.180] - Sam Harris

Obviously, you didn't pick your parents who didn't pick your genes. You didn't pick the environmental influences that determined every subsequent state of your brain. Right. You didn't create yourself. You won some sort of lottery there. But as will you point out, where you were born also was a major variable in your success. Very likely you didn't create the good luck not to be born in the middle of a civil war in a place like Congo or Syria or anywhere else, which would be hostile to, you know, many of the things you now take for granted.

[00:20:37.250] - Sam Harris
So there's something, frankly, obscene about not being sensitive to those disparities. And as you point out, living a good life and being the sort of person you are right to want to be has to entail some basic awareness of those facts and a compassionate impulse to make life better for people who are much less fortunate than we are. I mean, it's just if your vision of who you want to be doesn't include being connected to the rest of humanity and having compassion be part of the the operating system that oriented toward the shocking suffering of other people, even when it becomes proximate.

[00:21:28.940] - Sam Harris
You know, even when you're walking past Singer's Shallow Pond and you see someone drowning, you know, we have a we have a word for that orientation. And it's sociopathy or psychopathy. It's a false ethic to be so inured to the suffering of of other people that you can just decide to kind of close your accounts without even having to pay attention to it. And, you know, all under the rubric of being self-made. But, you know, none of this is to deny that in many cases things are better accomplished by business than by charity.

[00:22:05.330] - Sam Harris
Right. Or by government than by charity. So we're not denying any of that. I happen to think that building electric cars that people actually want to drive, you know, may be the biggest contribution to fighting climate change, or certainly one of them and maybe better than many environmental charities or managed to muster. I mean, so that there's a there are different levers to pull here to effect change in the world. But what also can't be denied is that there are cases where giving some of our resources to people or to causes that need them more than we do is the very essence of what it means to do good in the world that can't be disputed.

[00:22:43.790] - Sam Harris
And the Singer's Shallow Pond sharpens it up with a cartoon example. But it's really not such a cartoon when you think about the world we're living in and how much information we now have and how much agency we now have to affect the lives of other people. I mean, we're not isolated the way people were 200 years ago. And it is uncontroversial to say that anyone who would walk past a pond and decline to save a drowning child out of concern for his new shoes or his new suit, that person is a moral monster, and none of us want to be that sort of person.

[00:23:21.740]
And what's more, we're right to not want to be that sort of person.

[00:23:25.910]
But given our interconnectedness and given how much information we now have about the disparities in luck in this world, we have to recognize that though we're conditioned to act as though people at a distance from us, both in space and in time, matter less than people who are near at hand. If it was ever morally defensible, it's becoming less defensible because the distance is shrinking. We simply have too much information. So that's there's just so many. Ponds that are in view right now and a response to that is, I think, morally important, but in our last conversation, Will, you made a distinction that I think is very significant and it provides a much better framing for thinking about doing good.

[00:24:12.960]
And it was a distinction between obligation and opportunity. The obligation is Singer's shallow pond argument. You see a child drowning. You really do have a moral obligation to save that child or that there's just no way to maintain your sense that you're a good person if you don't. And then he forces us to recognize that really we stand in that same relation to many other causes, no matter how distant

we imagine them to be. But you favor the the opportunity framing of racing in to save children from a burning house.

[00:24:49.470]
Imagine how good you would feel doing that successfully. So let's just put that into play here, because I think it's a better way to think about this whole project. Yeah, exactly. So as I was suggesting earlier, just if I mean, for most people around the world, certainly in rich countries, if you look at your own values, well, one of those values is being a good person. And you can see this if you think about examples like.

[00:25:17.980]
There's you see a building on fire, there's a you know, a young girl kind of at the window and you kick the door down and you run in and you rescue that child. Like that moment would stay with you for the entire life.

[00:25:30.490]
You would reflect on that in your elderly years and think, wow, I actually really did something that was like that was pretty cool.

[00:25:37.100]
And it's just it's worth lingering there because everyone listening to us knows down to their toes that that would be, if not the defining moment in their life. You know, in the top five, there's just no way that wouldn't be one of the most satisfying experiences. You could live to be 150 years old, and that would still be in the top five most satisfying experiences of your life. And given what you're about to say, it's amazing to consider that and how opaque this is to most of us most of the time when we think about the opportunities to do good in the world.

[00:26:19.710]
Exactly. And yeah, I mean, continuing this, imagine if you did a similar thing kind of several times. So one week you saved someone from a burning building. The next week you saved someone from drowning. The month after that, you saw someone having a heart attack and you performed CPR and saved their life, too.

[00:26:37.840]
You'd think, wow, this is a really special life that I'm living.

[00:26:41.080]
But the truth is that we have that opportunity to be as much of a model hero, in fact, much more of a model hero every single year of our lives. And we can do that just by targeting our donations to the most effective charities to help those people who are poorest in the world. We could do that, too, if you wanted to choose a career that's going to have a really big impact on the lives of others.

[00:27:06.940]
And so it seems very unintuitive because we're in a very unusual place in the world.

[00:27:12.850]
You know, it's only over the last couple of hundred years such a wild discrepancy between rich countries and poor countries where people in rich countries have 100 times the income of the poorest people in the world and where we have the technology to be able to change the lives of people on other sides of the world, let alone the kind of technologies to, you know, imperil the entire future of the human race, such as the nuclear weapons or climate change. And so our moral instincts are just not attuned to that at all.

[00:27:44.200]
They are just not sensitive to the sheer scale of what an individual is able to achieve if he or she is trying to make a really positive difference in the world.

[00:27:54.940]
And so when we look at the, you know, history, look at the. Heroes like.

[00:28:00.520]
Think about William Wilberforce, aphasic Douglas, or the famous abolitionists, people who kind of campaigned for the end of slavery and the amount of good they did, all of these kind of great moral leaders and think, wow, these are really special people because of the amount they accomplished.

[00:28:15.160]
I actually think that's just attainable for many, many people around the world. Perhaps, you know, you're not quite going to be someone who can do as much as contribute to the abolition of slavery. But you are someone who can potentially save hundreds of thousands of lives or make a very significant difference to the entire course of the future to come. Hmm. Well, that's a great place to start. So now we will get into the details. OK, let's get into effective altruism.

[00:28:45.310]
Persay, how do you define it at this point? So the way I define effective altruism is that it's about using evidence and careful reasoning to try to figure out how to do as much good as possible and then taking action on that basis. And the real focus is on the most good. And that's so important because people don't appreciate just how great the difference and impact between different organizations are. When we've surveyed people, they seem to think that the best organizations are maybe 50 percent better than typical organizations like charities.

[00:29:22.510]
But that's not really the way of things. Instead, is that the best is more like hundreds or thousands of times better than a typical organization. And we just see this across the board when comparing charities, when comparing different sorts of actions. So for global health. You will save hundreds of times as many lives by focusing on antimalarial bed nets and distributing them than focusing on cancer treatment in the case of improving the lives of animals on factory farms will help thousands of times more animals by focusing on factory farms than if you try to help animals by focusing on pet shelters.

[00:29:57.680]
If you look at kind of risks to the future of civilization, manmade risks like novel pandemics are plausibly just, you know, a thousand times greater than magnitude than natural risks like, you know, asteroids that we might be more familiar with.

[00:30:12.680]
And that just means that focusing not just on doing some amount of good, but doing the very best is just it's so important. Because it's easy - yeah - it's easy just not to think about how wild this fact is. So, like, imagine if this were true of consumer goods. So at one store you want a beer. At one store, the beer costs 100 dollars another, it costs 10 cents.

[00:30:34.130]
That would just be completely mad. But that's the way things are in the world of trying to do good. It's it's like a ninety nine point nine percent off sale or 100000 percent exclusively by focusing on these best organizations, just the best deal you'll ever see in your life. And that's why it's so important for us to highlight this. OK, so I summarize effective altruism for myself now along these lines. So this this is a working definition, but it captures a few of the the areas of focus and the difference between solving problems with money and solving problems with your time or your choice of career.

[00:31:13.490]
In your response to my question, you illustrated a few different areas of focus so you could be talking about the poorest people in the world, but you could also be talking about long term risk to all of humanity. So the way I'm thinking about it now is that it's the question of using our time and our money to do one or more of the following things, to save the most number of lives, to reduce the most suffering, or to mitigate the worst risks of future death and suffering.

[00:31:43.550]
So then the question of effectiveness is, as you point out, there's so many different levels of competence and clarity around goals. There may be very effective charities that are targeting the wrong goals and they're ineffective charities targeting the right ones. And this does lend some credence to the the skepticism about charity itself that I referenced earlier. And there's one example here which does a lot of work in illustrating the problem. And this is something that you discuss in your book, Doing Good Better, which I recommend that people read.

[00:32:20.240]
But remind me about the ill fated play pump. Yes. So the now infamous play pump was a program that I got a lot of media coverage in the 2000s and even one the World Bank Development Marketplace Award. And the idea was identifying a the problem that many villages in sub-Saharan Africa do not have access to clean drinking water. And its idea was to install kind of children's merry go round one of the roundabouts, the things you purchase and then jump on and spin around, and that would harness the power of children's play in order to provide clean water for the world.

[00:33:06.420]
So the pushing on this merry go round. You would pump up water from the ground and it would act like a hand pump providing clean water for the village, and so people loved this idea. The media loved it, said, you know, providing clean water is child's play or that's the magic and they love to have fun on it.

[00:33:24.680]
So it was, you know, so they were hit. But the issue is that it was really a disastrous development intervention.

[00:33:31.940]
So none of the none of the local communities were consulted about whether they wanted the pump.

[00:33:36.770]
They liked their, you know, much cheaper, more productive, easier to use Zimbabwe and pumps that were sometimes, in fact, replaced by these pumps. And moreover, in fact, the pumps were sufficiently inefficient that one journalist estimated the children would have to play in the pump 25 hours per day in order to provide enough water for the local community. But obviously, children don't want to play on this matter grand old time, and so it would be left often to the elderly women of the village to push this brightly colored paper round and round.

[00:34:13.540]
One of the problems was that it didn't actually function like a merry go round where the you would gather momentum and keep spinning. It actually was just work to push, right? Well, exactly.

[00:34:25.000]
You need the point of a children's medical and as you push it and then you spin and if it's good, it's fairly well greased, it spins fairly. But you need to be providing energy into the system in order to pump water up from the ground. And so it wouldn't spin freely in the same way. It was enormous amounts of work.

[00:34:42.910]
Children would find it very tiring.

[00:34:44.140]
So it's just a fundamental misconception about engineering to deliver this pump in the first place. Yeah, absolutely.

[00:34:50.440]
And then there's just like, why would you think you can just go in and the place something that has

already been quite well optimized to the needs of the local people seems quite unlikely.

[00:35:01.450]
Like if this was such a good idea to ask the question, why wasn't it already invented? Why wasn't it already popular? There's not a compelling story about, well, it's a public good or something. There's a reason why you wouldn't have already been wouldn't have already been developed. And that's, you know, let alone the fact that the main issue in terms of water scarcity for people in the poorest countries is access to clean water and access to water.

[00:35:25.120]
And so instead, organisations like dispensers for safe water, which install chlorine like at the point of source. So at these hand pumps, clothing dispensers that they can easily put into the jerrycans that they use to carry water that sanitizes the water.

[00:35:41.380]
These are much more effective because that's really the issue, is dirty water rather than rather than access to water most of most of the time. OK, so this just functions as a clear example of the kinds of things that can happen when the story is better than the than the reality of of a charity. And if I recall correctly, there were you know, there were celebrities that got behind this and they raised had to be tens of millions of dollars for the play pump even after the fault in the very concept was revealed, they persisted.

[00:36:15.880]
I mean, there they kind of got locked in to this project. And I can't imagine it persists to this day. But they kept doubling down in the face of the obvious reasons to to abandon this project. This included, you know, kids getting injured on these things and, you know, kids having to be paid to run them. And it was a disaster anyway. You look at it. So this is the kind of thing that happens in various charitable enterprises.

[00:36:43.810]
And this is the kind of thing that if you're going to be effective as an as do you want to avoid. Yeah, absolutely. And just on the weather, they still continue. So I haven't checked in the last few years, but a few years ago when I did, they were still going and they were funded mainly by corporations like Colgate-Palmolive and obviously in a much diminished capacity because many of these failures were brought to light. And that was, you know, a good part of the story.

[00:37:11.020]
But what it does illustrate is a difference between the world of non-profits and the, you know, business world where in the business world, if you make a really bad product, then, well, at least if the market's functioning well, then the company will go out of business. You won't be able to sell it because the beneficiaries of the product are also the people paying for it. In the case of non-profits is very different. The beneficiaries are different from the people paying for the goods.

[00:37:35.800]
And so there's a disconnect between how well can you fund raise and how good is the program that you're implementing.

[00:37:41.680]
And so the sad fact is that bad charities don't die. Not nearly enough. Actually, that brings me to a question of about perverse incentives here that I do think animates sort of the more intelligent skepticism and is on precisely this point that, you know, charities, good and bad, can be incentivized to merely keep going. I mean, just just imagine a charity that solves its problem. It should be that, you know, if you're trying to, let's say, you know, eradicate malaria, you raise hundreds of millions of dollars to that end.

[00:38:19.480]
What happens to your charity when you actually eradicate malaria? We're not we're obviously not in

that position with respect to malaria, unfortunately. But there are many problems where you can see the charities are never incentivized to acknowledge that significant progress has been made and the progress is such that it calls into question whether this charity should exist for much longer. And, you know, there may be some, but I'm unaware of charities who are explicit about their aspiration to put themselves out of business because they're so effective.

[00:38:56.430]
Yes, I have a great example of this going along, so one charity I know of is called Scotts Care, and it was set up in the 17th century after the personal union of England and Scotland. And there were many Scots who migrated to London. And we were the poor. We were the indigent in London. And so it makes sense for there to be a non-profit helping make sure that poor Scots are, you know, had a livelihood with able to feed themselves and so on.

[00:39:27.990]
Is it the case in the 21st century? Poor Scots in London is the biggest global problem? No, it is not. Nonetheless, Scots care continues to this day over 300 years later.

[00:39:41.310]
Are there examples of charities that explicitly would want to put themselves out of business? I mean, giving what we can, which you joined is one I. Our ideal scenario is a situation where the idea that you would join a community because you're donating 10 percent is just weird, wild, like if you become vegetarian, very rare that you join the kind of a vegetarian society. Or if you decide not to be racist or decide not to be a liar, it's not like you join the no lion society or the no racist society.

[00:40:12.660]
And so that is what we're aiming for, is a world where it's just so utterly common sense that if you're born into a rich country, you should use a significant proportion of your resources to try and help other people impartially considered that the idea of needing a community or needing to be part of this kind of club without a group of people that just wouldn't even cross your mind.

[00:40:35.760]
So the day that they say that giving what we can is not needed is a very happy day for my perspective.

[00:40:42.810]
So let's talk about any misconceptions that people might have about effective altruism. Because, you know, the truth is, I've had some myself even having prepared to have conversations with with you and your colleague, Toby Ord. He's also been on the podcast. My first notion of effective altruism was that very much inspired by Peter Singer's shallow pond, that it really was just a matter of focusing on the poorest of the poor in the developing world, almost by definition.

[00:41:15.630]
And that's kind of the long and the short of it. And you're giving as much as you possibly can sacrifice. But the minimum bar would be 10 percent of your your income. What doesn't that capture about effective altruism? Yeah, thanks for bringing that up, because it is a challenge. We faced that the ideas that spread are the most mimetic with respect to effective optimism and not necessarily those that most accurately capture where the movement is, especially especially today.

[00:41:45.540]
So, as you say, many people think that effective altruism is just about earning as much money as possible to give to give well recommended global health and development charities.

[00:41:54.240]
But I think there's at least three ways in which that misconstrues things. One is the fact that there are just a wide variety of causes that we focus on now.

[00:42:03.840]
And in fact, among the kind of most engaged people in effective altruism, the biggest focus now is

making sure as future generations and making sure that things go well for the very many future generations to come, such as by focusing on existential risks that Toby talks about, like manmade pandemics, like, Heyy, animal welfare is another cause area.

[00:42:29.460]
It's not definitely by no means the majority focus, but is a significant minority focus as well. And this is lots of people trying to get better evidence and understanding of these issues and a variety of other issues, too. So voting the forum is something that I funded to an extent and champion to an extent. I'm really interested in more people working on the risk of war over the coming century. And then secondly, there are as well as donating, which is a very accessible and important way of doing it.

[00:43:01.770]
There's just a lot of in fact, the large majority of people within the effective altruism community are trying to make a difference, not primarily via their donations, though often they do donate to the primarily through their career choice by working in areas like research policy activism. And then just as a kind of framing in general, we just really don't think of effective altruism as a set of recommendations, but rather like a research project and methodology. So it's more like the science, you know, aspiring towards the scientific revolution than any particular theory.

[00:43:36.490]
Mm hmm. And what we're really trying to do is to do for the pursuit of goods what what the scientific revolution did for the pursuit of truth. It's an ambitious goal, but trying to make the pursuit of this more vigorous, more scientific enterprise. And for that reason, we don't see ourselves as kind of set of claims, but rather a living, breathing and evolving set of ideas. Yeah, yeah, I think it is useful to distinguish at least two levels here.

[00:44:05.390]
I mean, one is the specific question of whether an individual cause or an individual charity is a good one. And, you know, but by what metric would you even make that judgment? And how do we rank order our priorities? And like all of that is getting into the weeds of just what we should do with our resources. And obviously that has to be done. And I think the jury is very much out on on many of those questions.

[00:44:34.220]
And I want to get into those details going forward here. But the profound effect that your work has had on me thus far arrived at this other level of just the stark recognition that I want to do good in the world by default and I want to engineer my life such that that happens whether I'm inspired or not. The crucial distinction for me has been to see that there's the good feeling we get from philanthropy and doing good, and then there's the actual results in the world.

[00:45:13.100]
And those two things are only loosely coupled. This is one of the the worst things about us that we need to navigate around or at least be aware of as we live our lives, that we tend you know, we you know, human beings tend not to be the most disturbed by the most harmful things we do. And we tend not to be the most gratified by the most beneficial things we do. And we tend not to be the most frightened by the the most dangerous risks we run.

[00:45:43.070]
Right. And so it's just we're very easily distracted by good stories and and other bright, shiny objects. And the framing of a problem radically changes our perception of it. So the effect, you know, when you came on my podcast four years ago was for me to just realize, OK, well, now we're talking about give Wells most effective charities and the Against Malaria Foundation is at the top. I recognize in myself that I'm just not very excited about malaria or bed nets.

[00:46:15.560]
The problem isn't the sexiest for me. The remedy isn't the sexiest for me. And yet I rationally understand that if I want to save human lives, this is the dollar for dollar, the cheapest way to save a

human life. So that the epiphany for me is I just want to automate this and I'll just give every month to this charity without having to think about it. And so, you know, that is gratifying to me to some degree.

[00:46:45.860]
But the truth is, I almost never think about malaria or the Against Malaria Foundation or anything related to this project. And I'm doing the good anyway, because I just decided to not rely on my moral intuitions day to day and my desire to rid the world of of malaria. I just decided to automate it. The recognition that there's a difference between committing in a way that really takes it offline so that you no longer have to keep being your better self on that topic every day of the week.

[00:47:20.210]
It's just wiser and more effective to decide in your in your clearest moment of deliberation, you know, what you want to do and then just to build the structure to actually do that thing. And that's just one of several distinctions that, you know, you have brought into my understanding of how to do good. Yeah, absolutely.

[00:47:37.670]
I mean, it just we've got to recognize that we are these fallible, imperfect creatures. Donating is much like, you know, paying your pension or something, something you might think, oh, I really ought to do, but it's just hard to get motivated by. And so we need to exploit our own irrationality. And I think that comes in in two stages. First, like building up the initial motivation, you can sustain that for, you know, perhaps a feeling of moral outrage or just real kind of yearning to start to do something, you can get that.

[00:48:10.950]
So in my own case, when I was deciding how much should I try and commit to to give away over the course of my life, I just I looked up images of children suffering from horrific tropical diseases and that, you know, really stayed with me, kind of gave me that initial motivation or I still get that. If I read about the many close calls we had where we almost had a nuclear holocaust over the course of the 20th century, or if I, you know, learn more history and think about what the world would have been like if the Nazis had won the Second World War and created this global totalitarian state.

[00:48:44.850]
There's, you know, well, fiction like most recently reading 1984. And again, this kind of ways of just thinking just how bad and different the world could be that can really create the sense of like moral urgency or just, you know, on the news to the kind of moral outrages we see all the time. And then the second is how we don't how we director. And so in your own case, just saying, yes, every time I have a podcast, I donate three and a half thousand dollars and it saves a life.

[00:49:11.610]
Fairly good way of doing that. Similarly, you can have a system where every time a paycheck comes in, 10 percent of it just it doesn't even enter your bank account. It just goes to or at least immediately leaves to go to some. In fact, a charity that you've carefully thought about, and there's other hacks, too, so. Public commitments of a really big thing now, I think there's no way I'm backing out of my altruism now, right?

[00:49:35.410]
Too much of my identity is wrapped up in that now. So even if someone offered me, you know, a million pounds and I could skip town, you know, I wouldn't want to do it. I it's part of who I am. It's part of my social relationships. And that's, you know. Yeah, that's very that's very powerful to actually in a coming chapter here, I want to push back a little bit on how you are personally approaching given, because I think I have some rival intuitions here that I want to see how they survive contact with your sense of how you should live.

[00:50:09.580]
There's actually a kind of related point here where I'm wondering, when we think of causes that meet the test of the fact of altruism, they still seem to be weighted toward some obvious extremes. Right.

Like when you look at the value of a marginal dollar in sub-Saharan Africa or Bangladesh, you get so much more of a lift in human well-being for your money than you do or then you seem to. And in a place like the United States or the UK that, you know, by default, you generally have an argument for doing good elsewhere rather than locally.

[00:50:47.710]
But I'm wondering if this breaks down for a few reasons. So, I mean, just take an example like the problem of homelessness in San Francisco right now, leaving aside the fact that we don't seem to know what to do about homelessness, it appears to be a very hard problem to solve. You can't just build shelters for the mentally ill and substance abusers and call it a day, right? I mean, they quickly find that even they don't want to be in those shelters and, you know, they're back out on the streets.

[00:51:17.800]
And so you have to figure out what services you're going to provide. And there's all kinds of bad incentives and moral hazards here that then, you know, when you're the one city that does it, well, then then you're the city that's attracting the world's homeless. But let's just assume for the sake of argument that we knew how to spend money so that we could solve this problem. Would solving the problem of homelessness in San Francisco stand a chance of rising to the near the top of our priorities, in your view?

[00:51:46.600]
Yes.

[00:51:46.960]
So it would all just depend on how like the costs to save homelessness and how that compared with our other opportunities.

[00:51:54.850]
So in general, it's going to be the case that the very best opportunities are in order to improve.

[00:52:01.180]
Lives are going to be in the poorest countries because the very best ways of helping others have not yet been taken. So malaria is still life. It was wiped out in the U.S. and certainly by the early 20th century. It's an easy problem to solve. It's fairly cheap. And when we look at which countries, the problems that are still left are, you know, the comparatively harder ones to solve for whatever reason. So like in the case of homelessness, I'm not sure about the original source of this fact, but I have been told that.

[00:52:32.650]
So, yeah, for those who don't haven't ever lived in the Bay Area, the problem of homelessness is horrific, that there's just people with severe mental health issues, clear substance abuse, just like everywhere on the streets. It's so prevalent. It just amazes me that one of the richest countries in the world and one of the richest places within that country is unable to solve this problem. But I believe at least that in terms of funding at the local level, there's about 50000 dollars spent per homeless person in the Bay Area.

[00:53:04.240]
And what this suggests is that the problem is not to do with a lack of finances. And so if you were going to contribute more money, that it's unlikely to make an additional reason, perhaps at some perverse incentives. Fact, perhaps it's government bureaucracy, perhaps at some sort of legislation. I don't know. It's not an issue I know enough about.

[00:53:23.390]
But precisely because the U.S. is so rich, the San Francisco Bay Area, so rich, is that if this was something where we could turn money into a solution to the problem, it would more likely. More than likely. It probably would have happened already.

[00:53:37.090]
But that's not to say we'll never find issues in rich countries where you can do an enormous amount of good. So open philanthropy, which is kind of a core effectively to a foundation, one of its program areas is criminal justice reform that it started, I believe, about five years ago. And it really did think that the benefits to Americans that it could provide by funding changes to legislation to reduce the absurd rates of over incarceration in the U.S., where for context, the US incarcerates five times as many people as the UK does on a per person basis.

[00:54:17.590]
And there's a lot of evidence suggesting you could reduce that very significantly without changing rates of rates of crime.

[00:54:24.250]
It seems to be compatible to actually the best interventions in the poorest country.

[00:54:28.430]
Is, of course, this has now become an even more well focused issue, so I believe that they're finding it harder to now, you know, make a difference by funding organizations that wouldn't have otherwise be funded.

[00:54:39.980]
But this is at least one example where you can get things that come up that just for whatever reason, have not yet been funded. It's kind of new opportunities where you can do as much good. It's just that I think they're going to be comparatively much harder to find. He I think that this gets complicated for me when you look at just what we're going to target as a reduction in suffering. It is very easy to count dead people. Right.

[00:55:07.270]
So if we're just talking about saving lives, that's a pretty easy thing to calculate. If we can save more lives in Country X over country Y, well, then it seems like it's a net good to be spending our dollars in contracts. But when you think about human suffering and when you think about how so much of it is comparative, like the despair of being someone who has fallen through the cracks in a city like San Francisco could well be much worse.

[00:55:38.990]
I mean, I think there is certainly I don't know what data we have on this, but there's certainly a fair amount of anecdotal testimony that poor people in in a country like Bangladesh, while it's obviously terrible to be poor in Bangladesh, and there are many reasons to want to solve that problem. And by comparison, when you look at, you know, homeless people on the streets of San Francisco, they are not they're not nearly as poor as the poorest people in Bangladesh, of course.

[00:56:04.300]
And nor are they politically oppressed in the same way. I mean by global standards, they're barely oppressed at all. But it wouldn't surprise me if we could do a complete psychological evaluation or, you know, or just trade places with people in each condition. We would discover that the suffering of a person who is living in one of the richest cities in the world and is homeless and drug addicted and mentally ill or to, you know, pick off that menu of despair is actually, you know, the worst suffering on earth.

[00:56:41.570]
And again, we have to stipulate that we could solve this problem dollar for dollar in a way that, you know, we admit that we we don't know how to at the moment. It seems like just tracking the, you know, the GDP in each place and in the amount of money it would take to deliver a meal or get someone clothing or get someone shelter. And, you know, the power of the marginal dollar calculation doesn't necessarily capture the deeper facts of the case, or at least that's my concern.

[00:57:11.200]

So I'd actually agree with you on the question of, you know, take someone who. Yeah, the mentally unwell, they have drug addictions, the homeless in the San Francisco Bay Area. How bad is their day? And then take someone living in extreme poverty in India or sub-Saharan Africa. How bad is there a typical day? Yeah, I wouldn't want to make a claim that the. That homeless person in the US has a better life than the extreme poor, you know, I think it's it's not so hard to just hit rock bottom in terms of human suffering.

[00:57:47.750]
And I do just think that the homeless in the Bay Area just seem to have really terrible lives. And so the question, the question in terms of the difference of the how promising it is as a cause is much to do, much more to do with this question of whether the low hanging fruit has already been taken where, you know, just think about the most sick you've ever been and how horrible that was. And imagine, you know, and now think about that for months, having malaria, for example, and that that could you could have avoided that for a few dollars.

[00:58:24.560]
That's like, you know, an incredible fact. And that's where the real difference is, I think is in kind of the cost to solve a problem rather than necessarily like. The kind of poor person suffering, because while rich countries are in general happier than poorer countries, the worst off people, I mean, especially in the U.S., which has such a high variance in life outcomes. Yeah, the West, the lives of the West off people can easily be much the same.

[00:58:51.380]
Yeah, I guess there's some other concerns here that I have which and this speaks to a deeper problem with consequentialism, which is which is our orientation here, you know, not exclusively. And people can mean many things by that term. But there's just a problem in how you keep score because, you know, obviously there are bad things that can happen which have massive silver linings. Right. Which in the, you know, have good consequences in the end.

[00:59:18.530]
And there are apparently good things that happen that that actually have bad consequences elsewhere and or in the fullness of time. And it's hard to know when you can actually know that. You can assess what is what is true, the net, how you get to the bottom line of the consequences of any actions. But like when I think about the knock on effects of letting a place like San Francisco become a slum effectively. Right. Like you just think of like the exodus in tech from California at this moment, I don't know how deep or sustained it'll be, but I've lost count of the number of people in Silicon Valley who I've heard are leaving California at this point.

[01:00:02.240]
And the homelessness in San Francisco is very high on the list of reasons why that strikes me as a bad outcome that has far reaching significance for society. And again, the kind of thing that's not captured by just counting bodies or just looking at how cheap it is to buy bed nets, you know, and I'm sort of struggling to find a way of framing this that is fundamentally different from singers shallow pond that allows for some of the moral intuitions that I think many people have here, which is that there's an intrinsic good in in having a civilization that is producing the most abundance possible.

[01:00:47.330]
I mean, we want a highly technological, creative, beautiful civilization. We want gleaming cities with beautiful architecture. We want institutions that are massively well-funded, producing cures for diseases rather than just things like bed nets. Right. And we want beautiful art, the things we want. And I think there are things we we're right to want that are only compatible with the accumulation of wealth in certain respects. One, framing from singers, framing those intuitions are just wrong, or at least they're premature.

[01:01:30.390]
And we have to save the last child in the last pond before we can think about funding the Metropolitan Museum of Art right on some level. And many people are allergic to that intuition for reasons that I

understand. And I'm not sure that I can defeat Singer's argument here, but I have this image that essentially we have a lifeboat problem. Right. Like you and I are in the boat, we're safe. And then the question is, how many people can we pull in to the boat and save as well?

[01:02:01.940]
And, you know, as with any lifeboat, there's a problem of capacity. We can't save everyone all at once, but we can save many more people than we've saved thus far. But the thing is, we have a fancy lifeboat, right? A civilization itself is a is a fancy lifeboat. And there are people drowning and they're obviously drowning and we're saving some of them. And you and I are now arguing that we can save many, many more and we should save many, many more.

[01:02:29.540]
And anyone listening to us is lucky to be safely in this lifeboat with us. And the boat is not as crowded as it might be, but we do have finite resources in any moment. And the truth is, we're because it's a fancy lifeboat, you know, we are spending some of those resources on things other than reaching over the side and pulling in the next drowning person. So, you know, there's a bar that serves very good drinks and, you know, we've got a good Internet connection so we can stream movies.

[01:02:59.690]
And, you know, while this may seem perverse, again, if you extrapolate from here, you realize that I'm talking about civilization, which is a fancy lifeboat. And there's obviously an argument for spending a lot of time and a lot of money saving people and pulling them in. But I think there's also an argument for making the lifeboat better and better so that we have more smart, creative people incentivized to spend some time at the edge, pulling people in with better tools, tools that they only could have made.

[01:03:34.610]
Had they spent time elsewhere in the boat making those tools and this moves to the larger topic of just how we envision building a good society, even while there are there are moral emergencies right now somewhere that we we need to figure out how to respond to. Yeah, so this is crucially important set of questions, so the focus on knock on effects is very important. So when you again, let's just take the example of saving a life. You don't just save a life because that person goes on and does stuff.

[01:04:10.700]
Has they make the country richer? Perhaps they go and have kids. Perhaps, you know, they will emit CO2. That's a negative consequence. They'll innovate, they'll invent things. Maybe they got this huge stream basically from now till the end of time of consequences if you're doing this thing.

[01:04:28.130]
And it's quite plausible that the knock on effect, though, much harder to predict how much bigger effects than the short term effects, the benefits of the person you've whose life you saved or who you've benefited in the case of. Homelessness in the Bay Area versus extreme poverty in a poor country, I'd want to say that if we're looking at knock on effects of one, we want to do the same for both. So, you know, one thing I worry about over the course of the coming, you know, decades, but also even years is the possibility of a war between India and Pakistan.

[01:05:02.700]
But it's fact that rich democratic countries seem to not to go to war with each other. So one knock on effect of, you know, saving lives or helping development in India is perhaps we get to that point where India is of India is rich enough that it's not going to want to go to war because, you know, the cost benefit doesn't pay out in the same way. That would be another kind of potential good knock on effect. And that's not to say that the knock on effects favor the extreme poverty intervention compared to their homelessness is just that.

[01:05:31.290]
There's so many of them. It's very, very hard her to understand how these how these play out. And I think actually you then mentioned, well, we want to achieve some of the great things. So we want.

[01:05:45.400]
You know, to achieve the kind of highest apogee is of art, of development, I mean, a personal thing I'm sad that I will never get to see is the point in time where we just truly understand science, where we have actually figured out the fundamental laws, especially the fundamental physical laws, but also just great experiences to people having, you know, peaks of happiness that, you know, the very greatest achievements in of the present day just and the very greatest peaks of joy and ecstasy of a pleasant day just.

[01:06:20.440]
As basically almost, you know, insignificant in comparison, that's something that really I do think is important, but I think for all of those things, once you're then starting to take that seriously and take knock on effect seriously, that's the sort of reasoning that leads you to start thinking about what I call long term ism, which is the idea that the most important aspect of our actions is the impact we have over the very long run and will make us want to prioritize things like ensuring we don't have some truly massive catastrophe as a result of a nuclear war or a manmade pandemic that could derail this process of continued economic and technological growth that we seem to be undergoing or could make us want to avoid certain kind of just very bad value states like the lock in of a global totalitarian regime.

[01:07:11.410]
Another thing that I'm particularly worried about in terms of the future of humanity, or perhaps is just that we're worried the technological and economic growth will slow down.

[01:07:20.950]
And what we want to do is. Spur, you know, continued innovation into the future. And I think they actually are just really good arguments for that. I think I would be surprised if, though, if that is what your aim is, the best way of doing that goes via some things, such as focusing on homelessness in the Bay Area. Thousands find that kind of aim at those ends more directly. OK. Well, I think we're going to return to this concept of the fancy lifeboat at some point.

[01:07:52.720]
I do want to talk about your personal implementation of effective altruism in a subsequent lesson. But for the moment, let's get into the details of how we think about choosing a cause in the next chapter. OK, so how do we think about choosing specific causes? I've had my own adventures and misadventures with this since I took your pledge. Before we get into the specifics, I just want to point out a really wonderful effect on my psychology. That is I mean, you know, I've all I've always been, I think by real world standards, fairly charitable.

[01:08:32.800]
So giving to organizations that inspire me or who I think which I think are doing good work is not a foreign experience for me. But since connecting with you and now since taking the pledge, I'm now, you know, aggressively charitable. And what this has done to my brain is that there is a pure pleasure in doing this. And there's a kind of virtuous greed to help that gets kindled. And rather than seeing it as an obligation, it really feels like an opportunity.

[01:09:07.180]
I mean, just you want to run into that building and save the girl at the window. Absolutely. But across the street, there's a boy at the window and you want to you want to run in over there, too. And so it's this is actually a basis for psychological well-being. I mean, it makes me happy to put my attention in this direction. It's the antithesis of feeling like an onerous obligation. So anyway, I'm increasingly sensitive to causes that I that catch my eye and I want to support.

[01:09:35.230]
But I'm aware that I am a a malfunctioning robot with respect to my own, you know, moral compass. As I said, you know, I know that I'm not as excited about bed nets to stave off malaria as I should be and, you know, give in to that cause nonetheless, because I just recognize that the analysis is almost certainly sound there. But for me, what's interesting here is when I when I think about giving to a

cause that really doesn't quite meet the test, well, that then achieves the status for me of a kind of guilty pleasure.

[01:10:09.670]
Like, I feel a little guilty that I you know, I gave that much money to the homeless charity because, you know, Will just told me that that's not going to meet the test. So, OK, that's going to have to be above and beyond the 10 percent I pledged to the most effective charities. And so just having to differentiate the charitable donations that meet the test and those that don't is an interesting project. Psychologically, I don't know.

[01:10:33.040]
It's just it's just a very different territory than I've ever been with respect to philanthropy. But so this raises the issue.

[01:10:39.760]
So one of these charities is newly formed, right. Who does not yet have a long track record. I happen to know the people who are some of the people who created it. How could you fund a new organization with all these other established organizations that have track records that you can assess competing for your attention?

[01:10:58.660]
First thing I want to say is just does this count towards the pledge?

[01:11:02.170]
And one thing I definitely want to disabuse people of the notion of is that we think of ourselves as the authority of like what is effective. These are our best guesses we give well, or other organizations have put enormous amounts of research into this. But the still estimates there's plenty of things you can kind of disagree with. And it's actually quite exciting often to have someone come in and start disagreeing with us because maybe we're long and that's great. We can change our minds, have better beliefs.

[01:11:28.660]
And the second thing is that early stage charities absolutely can compete with charities with a more established that they could in just the same way as if you think about financial investment, you know, investing in bonds or the stock market is a way of making a return, but so is investing in startups. And if you had the view that you should never invest in startups, then that would definitely be a mistake. And actually quite a significant proportion of giffels expenditure each year on early stage nonprofits that have the potential in the future to become top recommended charities.

[01:12:02.110]
And so a set of questions I would ask for any organization I'm looking at is what's the cause that is focused on?

[01:12:08.800]
What's the program that's implementing and then who are the people who are who are kind of running that program?

[01:12:15.430]
But the kind of background is that there's just some things we know do enormous amounts of good and have this enormous amount of evidence for them. And so I feel like we want to be focusing on things where I do this like very promising evidence and we could. Potentially get more or something where in the nature of the beast, we cannot get very high quality evidence, but we have good, compelling arguments for thinking that this might be super important. So, you know, funding, clean energy, innovation, funding, you know, new developments in carbon capture and storage or nuclear power, something.

[01:12:47.990]

It's not like you can do a randomized controlled trial on that. But I think there's good kind of theoretical arguments for thinking that might be an extremely good way of combating climate change.

[01:12:56.870]
Is worth bearing in mind that like saying something that is the very best thing you can do with your money is an extremely high bar.

[01:13:02.420]
So, you know, if there's tens of thousands of possible organizations, they can only be one or two that, you know, have the they have the biggest bang for the buck. All right.

[01:13:12.710]
Well, it sounds like I'm opening a guilty pleasures fund to run alongside the Waking Up Foundation.

[01:13:18.080]
I'm very glad that the pleasures I'm glad that you are sufficiently you know, it's a very good instinct that you find out about these. Problems in the world, which are really bad and are motivated to want to help them. And so I'm really glad you think of them as pleasures. I don't think you should be beating yourself up, even if it doesn't seem like the very most optimal thing.

[01:13:40.040]
Yeah, yeah. No, no, I'm not. In fact, I haven't even guiltier pleasure to report which, you know, at the time I did it. You know, this is this is not through a charity. This is just a, you know, personal gift. But and this does connect back to just the kind of lives we want to live and how that informs this whole conversation. I remember I was listening to the The New York Times daily podcast, and this was when the covid pandemic was really peaking in the U.S. and everything seemed to be in freefall.

[01:14:13.760]
They profiled a couple who had a restaurant in, I think it was in New Orleans and they had an autistic child. And they were you know, everyone knows that restaurants were among the first businesses crushed by the pandemic for obvious reasons. And it was just a very affecting portrait of this family trying to figure out how they were going to survive and get their child the help. She I think it was a girl needed. So it was exactly the little girl fell down the well sort of story compared to the the genocide that no one can pay attention to because genocides are just boring.

[01:14:54.410]
And so I was completely aware of the dynamics of this. Helping these people could not survive comparison with just simply buying yet more bed nets. And yet the truth is, I really wanted to help these people. Right. So, you know, just sent the money out of the blue and it feels like an orientation that sort of two things here, that kind of rise to the defense of this kind of behavior. It feels like an orientation that I want to support in myself because it does seem like a truly virtuous source of mental pleasure.

[01:15:37.760]
I mean, it's better than almost anything else. I do the spending money selfishly and psychologically, but it's both born of a felt connection and it kind of ramify as that connection. And there's something about just honoring that bug in my moral hardware rather than merely avoiding it. That seems it seems like it is leading to just finding greater happiness in helping people in general, you know, in the most effective ways in, you know, middling, effective ways.

[01:16:08.990]
Feeling what I felt doing. That is part of why I'm talking to you now, trying to truly get my philanthropic house in order. Right. So it sort of seems all of a piece here. And I do think we we need to figure out how to leverage the salience of connection to other people and the pleasure of doing good. And if we lose sight of that, if we just keep saying that you can spend two thousand dollars here, which is better than spending three thousand dollars over there, completely disregarding the experience people are

having, engaging with the suffering of others, I feel like something is lost and I is another variable I would throw in here is, you know, this wasn't an example of this.

[01:16:52.320]
This wasn't a local problem. I was I was helping to solve. But had it been a local problem, had I been offered the opportunity to help my neighbor, you know, at greater than rational expense, that might have been the right thing to do. I mean, again, it's falling into the guilty pleasure been here compared to the absolutely optimized, most effective way of relieving suffering. But I don't know. I just feel like there's something lost if we're not in a position to honor a variable like locality ever.

[01:17:23.450]
We're not only building the world or affecting the world here. We're building our own minds. We're building the very basis by which we would continue to do good in the world in coming days and weeks and months and years.

[01:17:37.130]
Yeah. So, I mean, I essentially completely agree with you. I think it's really good that you supported that family. And yeah, it reminds me in my own case something that stayed with me. So I lived in Oakland, California for a while in a very poor, predominantly black neighborhood.

[01:17:53.660]
And I was just out. My man and woman kind of comes up to me. It's like I asks if I can stop and help for a second. And I thought she was going to want help, like carrying groceries or something. Be fine. It turns out she wanted me to move her couch, like, all the way down the street, took like two hours.

[01:18:12.740]
And I just don't. And that was my working day as well, because a lunch and I just don't like that the use of that time at all.

[01:18:21.110]
And why is that? And even from a rational perspective, I'm not saying that this is.

[01:18:26.370]
Oh, I shouldn't just merely shouldn't beat myself up or something, and I think it's because. Most of the time, we're just not the bigger question of like what individual action do we do, like in any particular case which. Kind of moral philosophy has typically focused on kind of act consequentialism, that's not typically the decisions we face. We face these much larger decisions like what career to pursue or something. Sometimes those are more like actions. But we also face the question of just what person to be, what kind of motivations and dispositions do I want to have.

[01:19:01.240]
And I think the idea of me becoming this like utility maximizing robot is like utterly cold and calculating at a time all the time, I think is certainly not possible for me, given just the fact that I am an embodied human being, but also probably not desirable either. I don't think you know, I don't think that an effective altruism movement would have started had we all been these cold utility maximizing robots. And so I think cultivating a personality such that you do get joy and award and motivation from being able to help people and get that feedback.

[01:19:37.420]
And that is like part of what you do in your life, I actually think can be the best way of living a life. When you consider your life as a whole and in particular, it's not necessarily doing those things, does not necessarily trade off very much at all and perhaps even help with the other things that you do. So in your case, you get this award from supporting this like poverty stricken family with a disabled child, or get the word from helping people in your local community that I'm assuming you can channel and like helps continue the motivation to do things that might seem much more alien or just harder to empathize with.

[01:20:15.400]
And I think that's OK. I think we should accept that. And that's, in fact, should be encouraged.

[01:20:19.690]
So, yeah, I think, like, it's very important once we take these ideas outside of the philosophy seminar room and actually try to live them to appreciate the instrumental benefits of doing these kind of everyday actions, as long as it ultimately helps you stand by this commitment to at least in part, to try and do just what we have, actually, all things considered. Think is going to be best for the world.

[01:20:42.850]
Yeah. So you mentioned that the variable of time here and this is another misconception about the effective of altruism, that it is only a matter of giving money to the most effective causes. You spent a lot of time thinking about how to prioritize one's time and think about doing good over the course of one's life based on how one's spends one's time. So in our next chapter, let's talk about how a person could think about having a career that helps the world.

[01:21:15.440]
OK, so we're going to speak more about the question of giving to various causes and how to do good in the world in terms of sharing the specific resource of money. But we're now talking about ones time. How do you think about time versus money here? And I know you've done a lot of work on the topic of how people can think about having rewarding careers that are net positive. And you have a website, 80000 hours that you might want to point people to hear.

[01:21:47.960]
So just let's talk about the variable of time and and how people can spend it to the benefit of others. So the organization is called 80000 Hours, because that's the typical number of hours that you work in the course of your life. If that's her, you know, approximately 40 year career working 40 hours a week, 50 weeks a year. So we use that to illustrate the fact that your choice of career is probably the altruistically speaking, the biggest decision you ever make.

[01:22:20.110]
It's absolutely enormous. Yet people spend very little of their time really thinking through. That question, I mean, you might think if you go out for dinner, then you spend maybe one percent of the time that you would spend at dinner thinking about where to eat like a few minutes or something, but spending one percent of 80000 hours on, you know, could a decision on what you should do, that would be 800 hours, enormous amount of time.

[01:22:44.600]
But I mean, why did I do philosophy? Well, I you know, I liked it at school.

[01:22:51.140]
I could have done maths, but my dad did maths and I wanted to differentiate myself from him. Like, I didn't have a very good reasoning process at all because we generally don't, you know, pay this nearly enough attention. And certainly when it comes to doing good, you have an enormous opportunity to have a huge impact through your career. And so what 80000 hours does via its website, via podcast and via a small amount of one on one advising is try to help people figure out which careers are such that they can have the biggest impact.

[01:23:23.370]
And in contrast, you know, this is a much you know, the question of what charities that I need to is exceptionally hard.

[01:23:30.260]
This is even harder, again, because firstly, you'll be working at many different organizations over the course of your life, probably not just one. And secondly, of course, there's a question of personal fact. Some people would be some people are good at some things and not others to do with them. And so how should you think about this? Well, the most important question, I think, is the question of what

caused to focus on and that involves big picture world view judgments and, you know, philosophical questions to so.

[01:24:01.190]
We tend to think of the question, of course, selection by using the heuristics of what causes and by a cause, I mean a big problem in the world, like climate change or gender inequality or poverty or factory farming or pandemics, the possibility of pandemics or I lock in the values. We look at those causes in terms of how important they are. That is how many individuals are affected by how much how neglected they are, which is how many resources are already going towards them.

[01:24:30.980]
And then finally, how tractable they are, how much we can make progress in this area. And in significant part because of those elastics. That's why we've effective altruism has chosen the focus areas it has, which includes pandemic preparedness, artificial intelligence, climate change, poverty, farm animal welfare, and potentially some others as well, like improving institutional decision making in some areas in scientific research. And so that's by far the biggest question, I think, because that really shapes the entire direction of your career.

[01:25:01.820]
And I think, you know, depending on the philosophical assumptions you put in can result in enormous, you know, differences and impact. Like, do you think animals count at all on or like a lot, I mean, would make an enormous difference in terms of whether you ought to be focusing on that simile? Like what weight do you give to future generations versus pleasant generations?

[01:25:23.060]
Potentially you can do hundreds of times as much good in one class area as you can in another. Yeah. And then within that, the question of where exactly to focus is going to just depend a lot on the particular cause area where different causes just have different bottlenecks.

[01:25:39.500]
We tend to find that, you know, working at the best nonprofits is often great. Research is often great, especially in kind of new, more nascent causes like safe development of artificial intelligence or epidemic pandemic preparedness. Often you need the such policy is often a very good thing to focus on as well. And in some areas, especially where, you know, money is the real bottleneck, then, you know, trying to do good through your donations primarily and therefore trying to take a job that's more lucrative can be the way to go to.

[01:26:11.300]
Yeah, that's a wrinkle that is kind of counterintuitive to people. The idea that the best way for you to contribute might in fact be to pursue the most lucrative career that you that you might be especially well placed to pursue. And it may have no obvious connection to doing good in the world, apart from the fact that you are now giving a lot of your resources to the most effective charities. So if you're a rock star or a professional soccer player or just doing something that you love to do and you have other reasons why you want to do it, but you're also making a lot of money that you can then give to great organizations, well, then it's hard to argue that your time would be better spent, you know, working in the nonprofit sector yourself or doing something where you wouldn't be laying claim to those kinds of resources.

[01:27:05.000]
Yeah, that's right.

[01:27:05.780]
And so it can be so within the effective altruism community, as is now, I think a minority of people are trying to do good in their career via the path of what's called earning to give. And again, it depends a lot on the course area, so.

[01:27:20.450]

What's the delay? You know, how much money is there relative to the size of the cause already? And, you know, in the case of things like scientific research or A.I. or pandemic preparedness, there's clearly just like a lot more demand for altruistically minded, sensible, competent people working in these fields than there is money. Whereas in the case of global health and development, there's just yeah, there were just these interventions and programs that we could scale up with hundreds of millions, billions of dollars that we just know work very well.

[01:27:52.140]
And they're that kind of money is kind of more of the bottleneck. And so kind of going back to these misconceptions about effective altruism, this idea of wanting to give, it's going again. It's fairly mimetic. People love the kinds of intuitive it is.

[01:28:06.660]
And, you know, is is one of the things we believe. But it's definitely kind of minority path, especially if you're focused on some of these areas where there already is a lot of potential funding. If you can, more about just how many people can we have working on these areas.

[01:28:22.470]
This raises another point where the whole culture around charity is not optimized for attracting the greatest talent. We have a double standard here, which many people are aware of. I think it's most clearly brought out by Dan Pelada. I don't know if you know him. He gave a TED talk on this topic and he organized some of the bike rides across America in support of various causes. I think the main one was AIDS. He might have organized a cancer one as well.

[01:28:53.370]
But, you know, these are ventures that raised, I think, hundreds of millions of dollars. And I think he was criticized for spending too much on overhead. But, you know, it's a choice where you can spend, you know, less than five percent on overhead and raise ten million dollars, or you could spend 30 percent on overhead and raise four hundred million dollars. What should you do? And it's pretty obvious you should do the latter if you're going to use those resources as well.

[01:29:18.990]
And yet there's a culture that prioritizes having the lowest possible overhead. And also there's this sense that if you're going to make millions of dollars personally by starting a software company or becoming a an actor in Hollywood or whatever it is, there's nothing wrong with that. But if you if you're making millions of dollars a year running a charity, well, then you're a greedy bastard. Right? And the idea that, you know, we wouldn't fault someone from pursuing a comparatively frivolous and even narcissistic career for getting rich in the meantime, but we would fault someone who's trying to cure cancer or save the most vulnerable people on Earth for getting rich while doing that.

[01:30:08.550]
That seems like a bizarre double standard with respect to how we want to incentivize people to marry. Because what what we're really demanding is someone come out of the most competitive school. And when faced with the choice of whether or not to work for a hedge fund or work for a charity doing good in the world, they have to also be someone who doesn't care about earning much money. So we need to sort of filtering for sainthood or something like sainthood among the most competent students at that stage.

[01:30:43.860]
And that seems less than optimal. I don't know how you view that. Yeah, I think it's I think it's a real shame so that, you know, newspapers have to publish rankings of the top paid charity CEOs and, you know, the God to just kind of scandal the charities, therefore ineffective, but.

[01:31:03.880]
What we should really care about, if we actually care about, you know, the potential beneficiaries, the people we're trying to help is just how much money are we giving this organization and how much good comes out the other end. And if it's the case that they can achieve more because they can

attract a more experienced and able person to lead the organization by paying more. Now, sure, that's like it's maybe a sad fact about the world. It would be nice if everyone were able to be maximally motivated purely by altruism.

[01:31:31.960]
But we know that's not the case. Then if they can achieve more by doing that, then, yeah, we should be encouraging them to do that.

[01:31:40.180]
You know, there's some arguments against like, oh, well, perhaps this kind of race to the bottom dynamics where if one organization starts paying more than other organizations should need to pay more to, and it just you get bloat in the system. I think that's the strongest case for the idea of low overheads when it comes to fundraising, because if one organization is fundraising, well, perhaps in part they're increasing the total amount of charitable giving that happens, but they're also probably taking money away from other organizations.

[01:32:13.030]
And so it can be the case that a general norm of lower overheads when it comes to fundraising is a good one. But when it comes to charity pay, we're obviously just radically far away from that and yeah, it shows that people are thinking about charity in a kind of fundamentally wrong way, at least, you know, for the effect of out his purposes we're thinking of, which is not thinking about it in terms of outcomes, but in terms of the virtues you demonstrate or how much you sacrifice or something.

[01:32:42.600]
And ultimately, when it comes to these problems that we're facing, these terrible injustices, this horrific suffering, I don't really care whether the person that helps is virtuous or not.

[01:32:51.240]
I just want the thing. I just want the suffering to stop. I just I just want people to be helped. And as long as they're not doing harm along the way, I don't think it really matters whether the people are paid a lot or a little more.

[01:33:04.260]
I think we should say something about the other side of this equation, which tends to get. Emphasized in most people's thinking about being good in the world, and this is the side of the consumer facing side of not contributing to the obvious harms in a way that is egregious or, you know, dialing down one's complicity in this unacceptable status quo as much as possible. And so this goes to things like becoming a vegetarian or a vegan or avoiding certain kinds of consumerism based on concern about climate change.

[01:33:41.730]
There's a long list of causes that people get committed to more. And in the spirit of negating certain bad behavior or polluting behavior, rather than focusing on on what they're, in fact, doing to solve problems or giving to specific organizations. Is there any general lesson to be drawn from the results of these efforts on both fronts? I mean, how much does harm avoidance as a consumer add to the scale of of merit here? What's the longest lever we can pull?

[01:34:16.360]
Personally, yeah. So I think there's a few things to say.

[01:34:19.320]
So right at the start I mentioned one of the key insights of effective altruism is this idea that different activities can vary by a factor of 100000 in terms of how much impact they have and even within ethical consumerism, I think that happens.

[01:34:35.370]
So if you want to cut out most animal suffering from your diet, I think you should cut out chicken and

pigs, maybe fish, whereas beef and milk, I think are comparatively small factors. If you want to reduce your carbon footprint, then giving up beef and lamb, reducing transatlantic flights that you're seeing, how much you drive makes significant differences and dozens of times as much impact as things like the cycling or upgrading light bulbs or using plastic bags from the purely consequentialist outcome based perspective.

[01:35:07.230]
I think it is systematically the case that these ethical consumers and behaviors are small in terms of their impact compared to the impact that you can do via your donations of your career.

[01:35:19.710]
And the reason is just there's a fairly limited range of things that you can do by changing your consumption behavior. There's just things you are buying anyway, and then you can stop. Whereas if you're donating or you're choosing a career, then you can choose the fairly most effective things to be doing. So take the case of being vegetarian. So I've been vegetarian for 15 years now. I have no plans of stopping that.

[01:35:43.920]
But if I think about how many animals I'm helping in the course of a year as a result of being vegetarian, and how does that compare when I'm looking at the effectiveness of the very most effective animal welfare charities?

[01:35:57.200]
Which is typically what are called kind of corporate campaigns. So it turns out the most effective way of reducing the number of that we know of reducing the number of hens and factory farms, laying eggs in just the most atrocious, terrible conditions of suffering seems to be by like campaigning large retailers to change the eggs they purchased in the supply chain. You can actually get almost a lot of push there. And the figures are just like astonishing. It's like something like, you know, 50 animals that you're preventing the significant torture of for every dollar that you're spending on these campaigns.

[01:36:34.640]
And so if you just do the maths, like the amount of good you do by becoming vegetarian is equivalent to the amount of good you do by donating a few dollars to these very most effective campaigns. I think similarly is true for the juicing your carbon footprint. My current favorite climate change charity, Clean Air Task Force, which lobbies the U.S. government to improve its regulations around fossil fuels and promotes energy innovation as well, I think probably reduces a ton of CO2 for about about a dollar.

[01:37:06.170]
And that means if you're in the U.S., an average U.S. citizen emits about 16 tons of carbon dioxide equivalent.

[01:37:12.380]
If you did all of these most effective things of cutting out meat and all your transatlantic flights and getting rid of your car and so on, you might be able to reduce that sextons or so. And that's, you know, the same as giving about six dollars to the most effective charities.

[01:37:28.640]
And so it just does seem that these are just much more powerful from the perspective of outcomes. The next question philosophically is whether you have some non consequentialist reason to do these things. And they have I think it differs. So I think the case is much longer for becoming vegetarian than for climate change, because if I buy a factory farm chicken and then donate to a corporate campaign, well, I probably harmed different chickens.

[01:38:00.320]
And it seems like that's you know, you can offset the harm to one individual by a benefit to another individual. Whereas if I have a lifetime of emissions, but at the same time donate a sufficient amount

to climate change charities, I've probably just reduced the total amount of CO2 going into the atmosphere over the course of my lifetime.

[01:38:19.520]
And there isn't anyone who's harmed an expectation, at least by the entire course of my life. And so it's not like I'm trading a harm to one person for the benefit to another. But these are quite these are quite subtle issues when we get into these kind of non consequentialist reasons.

[01:38:37.040]
Yeah, there are also ways in which the business community and innovation in general can come to the rescue here. So, for instance, there's a company, I believe the name is going to be changed. But it was it was called Memphis' Meets that is spearheading this this revolution in what's called cultured meat or clean meat, where they, you know, they take a single cell from an animal and amplify it. So no animals are killed in the process of making these steaks or these meatballs or these chicken cutlets.

[01:39:07.460]
And they're trying to bring this to scale. And I had the CEO whom a on my podcast a couple of years ago and actually invested in in the company along with many other people. And hopefully this will bear fruit. That set an example of something where though it was unthinkable some years ago, we might suddenly find ourselves living in a world where you can buy steak and hamburger meat and pork and chicken without harming any animals. And it may also have other significant benefits like cutting down on zino viruses.

[01:39:43.160]
And and, you know, that connects to the pandemic risk issue. I mean, we're really working our factory farms are wet markets of another sort. And so it is with climate change on some level, we're waiting and expecting for technology to come to the rescue here where it's bringing down the cost of renewable energy to the point where it's there is literally no reason to be using fossil fuels or bringing us a new generation, a new generation of nuclear reactors that don't have any of the downsides of old ones.

[01:40:15.770]
And again, this does connect to the concern I had around the fancy lifeboat. We have to do the necessary things in our lifeboat that allow for those kinds of breakthroughs because, you know, those are the in many cases the solutions that just fundamentally take away the problem rather than merely mitigate it. Yes.

[01:40:38.980]
So I totally agree. And I think that. So in the case of. You know, if you're trying to alleviate animal suffering by as much as possible, I think that, yeah, funding research into clean meats, plausibly the best thing you can do, it's hard to make a comparison with the more direct campaigns, but definitely possibly the best in the case of climate change have recently been pretty convinced that the most effective thing we can be doing is promoting clean energy innovation.

[01:41:07.310]
In this case, this is another example of importance versus neglecting this where you mentioned renewables and they are really key part of the part of the solution that other areas really notably more neglected. So carbon capture and storage where you're capturing CO2, is it as it emerges from fossil fuel, power plants and nuclear power get quite a small amount of funding compared to solar and wind, even though the Intergovernmental Panel on Climate Change thinks that they're also a very large part of the solution.

[01:41:40.010]
But here, I think the distinction is focusing on issues in which countries in order to benefit people in those rich countries or kind of as a means to some other sort of benefit. And so I think it's very often the case that you should focus on, like you might be sending money towards things happening in a rich country like the U.S., but not because you're trying to benefit people in the U.S., because you're

trying to benefit the world. So maybe you're funding.

[01:42:05.300]
Yeah, there's clean meat start up or you're funding research into low carbon forms of energy. And sure, like that might happen in the US, which is still the world's research leader. That's very justified. But that's kind of partly the beneficiaries and the use of these things. But it's also it's you know, it's global. It's future generations, too. You're kind of influencing, as it were, the people who are in positions of power, who have the most influence over how things are going to go into the future.

[01:42:36.180]
OK, so in our next chapter, let's talk about how we build effective altruism into our lives and just make this as personally actionable for people as we can. OK, so we've sketched the basic framework of effective altruism and just how we think about systematically evaluating various causes, how we how we think about, you know, what would be prioritized with respect to things like actual outcomes versus a good story. And we've referenced a few things that are now in the the effect of altruist canon, like giving a minimum of 10 percent of one's income a year.

[01:43:22.100]
And that's really, if I'm not mistaken, you just took that as a nice round number that people had some traditional associations with. You know, in religious communities, there's a notion of tithing that amount. And it seemed like not so large as to be impossible to contemplate, but not so small as to be ineffectual. Mm hmm. Maybe let's start there. How so? Am I right in thinking that the 10 percent number just was kind of pulled out of a hat, but seemed like a good starting point, but there's nothing about it that's carved in stone from your point of view.

[01:43:57.020]
Exactly. It's not it's not a magic number, but it's just it's in this Goldilocks zone where Toby originally had had the thought that he would be promoting what he calls the further pledge, which is where you just set a cap on your income and give everything above that. But the issue, I think, seems pretty clear that if he'd been promoting that, well, very few people would have joined him. We do have a number of people who've taken the further pledge, but it's a very small minority of the 5000 members we have.

[01:44:26.150]
On the other hand, if we were promoting a one percent pledge, let's say, well, we're probably just not changing people's behavior compared to how much they donate anyway.

[01:44:34.070]
So in the UK, people donate on average point seven percent of the income. And the US, if you include educational donations and church donations, people donate about two percent of their income. So if I was saying, oh, we should donate one percent, probably those people would have been giving one percent anyway. And so we thought 10 percent is in this Goldilocks zone. And like you say, it has this long history where for religious reasons, people much poorer than us.

[01:45:00.200]
And earlier, historically parks have, you know, been able to donate 10 percent. We also have 10 fingers. It's a nice round number.

[01:45:08.030]
But, you know, many people who are part of the community donate much more than that. Many people who, you know, affirm call people part of the effective altruism community don't donate that much. They do good via other other ways. Instead, it's interesting to consider the psychology of this because I can imagine many people entertaining the prospect of giving 10 percent of their money away and feeling, well, I could easily do that if I were rich, but I can't do that now.

[01:45:39.320]
And I can imagine many rich people thinking, well, that's. A lot of money, right? It's like every year after year I'm making a lot of money, and you're telling me year after year after year I'm going to give 10

percent away, you know, that's millions of dollars a year. So it could be the fact that there's no point on the continuum of of earning where if you're of a certain frame of mind, it's going to seem like a Goldilocks value isn't you either feel too poor or too rich and there's no sweet spot or, you know, to flip that around, you can recognize that however much money you're making, you can always give 10 percent to the most effective ways of alleviating suffering.

[01:46:25.640]
Once you have this epiphany, you can always find those 10 percent at every point. And if you're not making much money, obviously 10 percent will be a small amount of money. And if you're making a lot of money, it'll be a large amount. But there's almost all it's almost always the case that there's 10 percent of fat there to be found. So, yeah. So did you have thought about just the psychology of someone who feels not immediately comfortable with the idea of making such a commitment?

[01:46:53.660]
Yeah, I think there's two things I'd like to say to that person. One is the kind of direct argument. The second is more pragmatic. The direct one is just that even if you feel like, oh, I could donate that amount if I were rich, probably you are rich. And if you're listening to this, so as a if you're single and you earn sixty six thousand dollars, then you're in the global one percent of the world in terms of income distribution.

[01:47:21.680]
And what's more, even after donating 10 percent of your income, you would still be in the richest one percent of the world's population. If you earn thirty thirty five thousand dollars, which, you know, we would not think of as being a rich person, even after donating 10 percent, still being the richest five percent of the world's population and learning those facts was very motivating for me when I first started thinking about my giving.

[01:47:45.080]
So that's kind of a direct argument for the more pragmatic one is to think, well, if you're at most stages in your life, you'll be earning more in the future than you are now.

[01:47:57.110]
You know, people's incomes tend to increase over time, and you might just reflect, well, how do I feel about money at the moment? And if you feel kind of ovie about it, you know, perhaps you're in a situation where you're like, oh, no, I'm actually just very worried. It's like serious health issues or something. Then it's like, OK, we'll take care of that first, if you like. Well, actually, you know, most people, they don't think additional money will make that much of a difference.

[01:48:21.020]
Then what you can do is just think, OK, maybe I'm not going to give up to 10 percent now, but I'll give a very significant proportion of the additional money I make any future raises.

[01:48:31.650]
So maybe I get 50 percent of that amount and probably after that means that you're still increasing the amount you're earning over time. But at the same time, you're you know, if you do that, then over the few years you're probably quite soon end up giving 10 percent of your overall income. So at no point in this plan do you ever have to go backwards, as it were, living on less. In fact, you're always earning more, but yet you're giving more at the same time.

[01:48:55.850]
And I've certainly found that in my own life where, you know, I started thinking about giving as a graduate student. So, you know, I now earn you know, I now live on like twice as much more than twice as much as I did when I first started giving. But I'm also able to give, you know, a significant amount of of my income.

[01:49:15.080]
Remind me, how have you approached this personally? Because you haven't taken a minimum 10

percent plans. You think of it differently. So what have you done over the years? Yeah. So, you know, so I have taken the giving what we can pledge, which is 10 percent of any point. And then I also have intention and and plan to donate everything above what is the equivalent of 20000 pounds per year in Oxford 2009, which is now about twenty seven thousand pounds per year.

[01:49:44.060]
I've never written this down. It's like a formal pledge. The reason being that they were just too many possible kind of exceptions. So if I had kids, I'd want to increase that. If there were situations where I thought my ability to do good in the world would be like very severely hindered, I'd want to kind of avoid that. But that is the amount that I'm giving at the moment and it's the amount I plan to give for the rest of my life, just so I understand.

[01:50:06.900]
So you see, you're giving anything you make above twenty seven thousand pounds a year to charity.

[01:50:14.400]
That's right. Yeah, that's right. Post-tax. And so my income is a little bit complicated in terms of how you evaluate it, because it's my university income, but then also book sales and so on. I think on the most natural and it's things like speaking engagements. I don't take that. I could I think on the most natural way of doing it, I give a little over 50 percent of my income.

[01:50:36.020]
So I want to explore that with you a little bit, because, again, I'm returning to our fancy lifeboat and wondering. Just how fancy it can be in a way that's compatible with the project of doing the most good in the world and what I detect in myself and in most of the people I meet, and I'm sure in this is an intuition that is shared by many of our listeners. Many people will be reluctant to give up on the aspiration to be wealthy with everything that that implies.

[01:51:11.130]
You know, obviously they want to work hard and make their money in a way that is good for the world, or at least benign. They can follow all of the ethical arguments that would say, you know. Right. Livelihood in some sense is important. But if people really start to succeed in life, I think there's something that will strike many people, if not most, as to abstemious and monkish about the lifestyle you're advertising in choosing to live on that amount of money and give away everything above it or even just, you know, giving away 50 percent of one's income.

[01:51:56.400]
And again, I think it does actually connect with the question of effectiveness. I mean, like it's at least possible that you would be more effective if you were wealthy and living with all that, all that that entails living as a wealthy person. And I mean, just to take by example someone like Bill Gates, you know, he's obviously the most extreme example I could find because he's you know, he's one of the wealthiest people on earth. Still, I think is number two, perhaps.

[01:52:26.160]
And he's also the the probably well established now. He's the biggest benefactor of charity in human history. Perhaps the Gates Foundation has been funded to the tune of tens of billions of dollars by him at this point. And so I'm sure he's spent a ton of money on himself and his family. Right. His his life is probably filled to the brim with luxury, but his indulgence in luxury is still just a rounding error on the amount of money he's giving away.

[01:52:58.290]
Right. So it's it's actually hard to run a counterfactual here, but I'd be willing to bet that Gates would be less effective and less wealthy and and have less money to give away if he were living like a monk in any sense. And I think maybe more importantly, his life would be less inspiring, less inspiring example to many other wealthy people. If Bill Gates came out of the closet and said, listen, I'm I'm living on fifty thousand dollars a year and giving all my money away to charity, that wouldn't have the same kind of kindling effect, I think is life at this point is in fact having, which is you can really have your cake and

eat it, too.

[01:53:41.010]
You can be a billionaire who lives in a massive smart house with all that sexy technology, even fly around on a private jet and be the the most charitable person in human history. And if you just think of the value of his time. Right, like if he were living a more abstemious life and I mean, just imagine the side of Bill Gates spending an hour trying to save 50 dollars on a on a new toaster oven. Right. You know, bargain hunting.

[01:54:10.350]
It would be such a colossal waste of his time given the value of his time. Again, I don't I don't have any specifics, really, about how how to think about this counterfactual. But I do have a general sense that it actually this is actually a point you made in our first conversation, I believe, which is you don't want to be an anti-hero in any sense. Right. You want to like if you can inspire only one other person to give at the level that you are giving, you have doubled the good you can do in the world.

[01:54:40.350]
So on some level, you want your life to be the most compelling advertisement for this whole project. And I'm just wondering if I mean, for instance, I'm just wondering what changes we would want to make to, you know, Bill Gates, his life at this point to make him an even more inspiring advertisement for effective altruism to other very, very wealthy people. Right. And I mean, it might be dialing down certain things, but given how much good he's able to do him, buying a fancy car is just it doesn't even register in terms of actual allocation of resources.

[01:55:21.540]
So anyway, I pitched that to you. Yeah, it's terrific. I think so.

[01:55:25.800]
There's three different stands I think I'd like to pick apart. So the first is whether everyone should be like me. And I really don't want to make the claim. I certainly don't want to say, well, I can do this thing so everyone else can because I really just think I am in a position of such. ATAP Ravlich, so being, you know, born into, you know, a middle class family in a rich country, being privately educated, going to Cambridge, then Oxford and Cambridge and Oxford, being like like tall and male and white and broadly they like and then also just having kind of inexpensive tastes like my ideal day involves sitting on a couch and drinking tea and reading some interesting new research and perhaps like.

[01:56:15.550]
Doing going well, swimming it all. Yeah, and then secondly, also, I have just these amazing benefits in virtue of the work that I do. I have this incredibly like I meet these incredibly varied, interesting kind of array of people. And so I just don't really think I could stand here and say, well, everyone should do the same as me, because I think I've just had it kind of so easy that it doesn't really feel like, you know, if I think about the sacrifices I have made or the things I found hard over the course of 10 years.

[01:56:48.350]
That's much more like doing scary things, like being on the Sam Harris podcast or doing a TED talk or, you know, meeting, you know, fairly wealthy or very important people, things that might kind of cause anxiety much more than the kind of financial side of things. But I recognize there are other people for whom, like money just really matters to them. And I think you just in part, you're kind of born with a set of preferences on these things, or perhaps they're molded early on in childhood and you don't necessarily have control over them.

[01:57:18.830]
So that's kind of me as an yeah. What I'm trying to convey with this second is the time value of money. And this is something I've really wrestled with because it just is the case that in terms of my personal impact, my donations are just. A very small part of that, because, you know, we have been successful, we are, you know, giving what we can, has now moved 200 million dollars. This about one and a half billion dollars of pledged donations.

[01:57:49.950]
The movement as a whole certainly has over 10 billion dollars of assets that kind of will be going out. And then, you know, I'm donating my, you know, thousands of pounds. Yeah. And that does not make tens of thousands of pounds per year. And it's is just very clearly kind of small on the scale. And so that's definitely something I've wrestled with. I don't think I lose enormous amounts of time. My guess is that it's maybe a couple of days of time a year.

[01:58:16.510]
I have done some things. So like, you know, via my work, I have an assistant. If I'm doing business trips like that counts as expenses and my personal money. So I'm trying to keep it separate. There's some things you can't do. So like if you live close to your office, you know, I can't count that as a business expense, but shorten your commute. So it's not, like, perfect as a way of doing that.

[01:58:38.290]
And so I do think there's an argument, an argument against that. And I think that is definitely a reason of caution for making kind of a very large commitment. And then the final aspect is. Yeah, what sort of message you want to send. And probably my guess is that you just want a bit of market segmentation here where some people should you know, some people should perhaps show what can be done. Others should show. Well, no, actually, you can have this amazing life while, you know, not having to wear the hair shirt and so on.

[01:59:09.940]
You know, I think perhaps you could actually convince me that maybe I'm, you know, sending the wrong message and would do more good if I had some other sort of pledge. And maybe you would be right about that. I definitely when I made these plans, I wasn't thinking through these things quite as carefully as I was now. But I did want to I just want to just kind of show a proof of concept. Yeah. I guess I'm I'm wondering if there's a path through this wilderness that doesn't stigmatize wealth at all.

[01:59:39.450]
In the end game for me, in the presence of absolute abundance is, you know, everyone gets to live like Bill Gates on some level. If we make it if we get to the twenty second century and we solve the alignment problem and now we're just pulling wealth out of the ether, I mean, essentially just we've got Deutsch's universal constructor's, you know, building every machine atom by atom, and we can do more or less anything we want.

[02:00:08.530]
Well, then this can't be based on an ethic where wealth is. Is it all stigmatized or should have opprobrium attached to it is a total disconnection from the suffering of other people and comfort with the more shocking disparities in wealth that we see all around us. Once a know, reasonably successful person signs on to the the effect of altruist ethic and begins thinking about his or her life in terms of earning to give. On some level, there's a flywheel effect here where one's desire to be wealthy actually amplifies one's commitment to giving.

[02:00:56.140]
So that in part the reason why you would continue working is because you have an opportunity to give so much money away and do so much good and it kind of purifies one's earning in the first place. I mean, like I can imagine, most wealthy people get to a point where they're making enough money so that they don't have to worry about money anymore. And then there's this question, well, why am I making all this money? Why am I still working?

[02:01:24.250]
And the moment they decide to give a a certain amount of money away a year, just algorithmically, then they feel like, OK, if this number keeps going up, that is a good thing. Right? So, like, I can get out of bed in the morning and know that today, you know, if it's 10 percent, you know, one day and 10 is given over wholly to solving the worst suffering or saving the most lives or mitigating the the worst long term risk.

[02:01:51.340]
And if it's 20 percent, it's, you know, two days out of ten. And if it's 30 percent, it's three days out of ten. And they could even dial it up. I'm I'm just imagining, let's say somebody is making ten million dollars a year and he thinks, OK, I can sign on and give 10 percent of my income away to charity. That sounds like the right thing to do. And he's persuaded that this should be the minimum.

[02:02:11.260]
And he but he then aspires to scale this up as he earns more money. You know, maybe this would be the algorithm. You know, for each million he makes more a year. He just adds the percentage of he's making. Fourteen million one year, he'll he'll give fourteen percent of his income away. If it's 50 million, he'll give 50 percent away. Right. And obviously, if let's say the minimum he wants to make is nine million a year, well, then he can get to.

[02:02:37.320]
Up to 91 percent of 100 million dollars a year, you can give that away, but I can imagine being a very wealthy person who, you know, as you're scaling one of these outlier careers, it would be, you know, fairly thrilling to be the person who's making 100 million dollars that year, knowing that you're going to give ninety one percent of that away to the most effective charities. And you might not be the person who would have seen any other logic in driving to that kind of wealth.

[02:03:10.890]
You know, when you were the person who was making ten million dollars a year because ten million dollars a year was good enough, I mean, obviously you can live on that. You know, there's nothing that materially is going to change for you as you make more money. But because he or she plugged in to this this earning to give logic and in some ways the greater commitment to earning was was leveraged by a desire to maintain a wealthy lifestyle.

[02:03:37.530]
Right. So I guess this person does want nine million dollars a year. Right, every year. But now they're much wealthier than that and giving much more money away. I'm just trying to figure out how we can capture the imagination of people who would see the example of Bill Gates and say, OK, that's that's the sweet spot, as opposed to any kind of example that, however, subtly stigmatizes being wealthy in the first place. Hmm. Yeah, I think these are good points and.

[02:04:09.900]
It's true, I think the stigma and wealth, per say, is not a good thing where, you know, if you build a company that's doing good stuff and people like the products and they get value from it. And so there's enormous surpluses of gains from trade and you get wealthy as a result of that. That's a good thing. Obviously, there are some people who, like, make enormous amounts of money doing bad things, selling opioids or building factory farms, but I don't think that's the majority.

[02:04:37.670]
And I do think it's the case that, you know, it's kind of like optimal taxation theory. But you the weird thing is that you're imposing the tax on yourself where depending on your psychology, if you, you know, say I'm going to give 100 percent as the highest tax rate. Well, you're not incentivized to earn anymore. And so the precise amount that you want to give is just quite sensitive to this question of just how unmotivated you're going to be in order to earn more.

[02:05:05.720]
So in my own case, you know, I'm not it's very clear that I'm the way I'm going to do good is not climatically via my donations. So perhaps this disincentive effect is not very important. But if my aim were to get as little as possible, then, well, I need to really look inside my own psychology, figured out how much, especially over the entire course of my life, can I be motivated by pure altruism versus self-interest.

[02:05:33.230]

And I strongly doubt that the kind of optimal tax rate would be, you know, via my donations would be 100 percent. It would be something in between.

[02:05:41.540]
That's what I'm kind of fishing for here. And, you know, by no means am convinced that I'm right. But I'm just wondering if, in addition to all the other things you want, you know, as revealed in this conversation, you know, for yourself and the world and, you know, acknowledging that your you know, your primary contribution to doing good in the world might, in fact, be your ideas and your ability to get them out there.

[02:06:05.720]
I mean, like you've had the effect you've had on me, and I'm going to have my effect on on my audience. And, you know, conversations like this have the effect that they have. And so there's no question you are inspiring people to marshal their resources in these directions and think more clearly about these issues. But what if it were also the case that if you secretly really wanted to own a Ferrari, you would actually make different decisions such that in addition to all the messaging, you would also become a very wealthy person, giving away a lot of money?

[02:06:42.470]
Yeah, I mean, if it was the case, you know, if it was the case that I was planning to earn, to give. And so I think a very common kind of figure for people who are going to earn, to give via entrepreneurship or other high, high end and careers is a 50 percent figure where they plan to give half of what they earn at least once they start earning a significant amount. And that has seemed to work pretty well from the people I know.

[02:07:08.990]
It's also notably the figure that Bill Gates uses for his giving pledge where billionaires can join the giving pledge if they give at least 50 percent of their income of their wealth. Most take that pledge, if I'm not mistaken, is pushed off to the end of their life, rather, just imagine they're going to give it upon their death to charity.

[02:07:29.330]
So you are allowed to do that? I know. I don't know exactly the proportions. It varies like the tech. Tech founders tend to give earlier than other sorts of people, I'm also not just I'm actually a little bit confused about what pledging 50 percent of your wealth means.

[02:07:45.470]
So if I am a billionaire one year and then lose half my money and I've got 500 million dollars the next year, do I have to give half of that or do I have to give half of the amount? When I pledge, which would have been all my money anyway, confuses me a little bit. The details of it. But it is the case that yeah, you can fulfill your pledge completely in the giving pledge by donating entirely after your death.

[02:08:11.120]
And there are questions about how much people actually. Fulfill these pledges to right, but then yeah, and I think and I really do want to say, like, that's also just quite reasonable. Different people have different attitudes to money. I think it's a very rare person indeed that can be entirely motivated, you know, and because we're talking about motivation over decades and we're talking about every single day, motivation can be motivated at all times by pure altruism.

[02:08:38.030]
I think that's very hard. And so if someone instead wants to pick, you know, percentage number and to that, that seems like a sensible way to go. And in particular, you want to be sustainable where if it's the case that moving from, I don't know, 50 percent to 60 percent means that actually you like your desire to do all of this kind of burns out and you go and do something else that's very bad indeed. And you want to be someone who's like.

[02:09:06.090]

You know, I think the right attitude you want to have towards giving is not to be someone where it's like, oh yeah, I'm giving this amount, but it's just so hard. And I like I really don't like my life and it's unpleasant. That is, you know, not an inspiring message. Julia Wise has this wonderful member of the effects of Elton's community, has this wonderful post called cheerfully, where she talks about having kids and thinking about that as a question and says that know what you want to be is this model, this ideal where you are doing what you're doing and you're saying, God, my life is great, I'm able to do this and I'm still having a really wonderful life.

[02:09:41.700]
That's certainly how I feel about my life. And I think for many people who are going into these higher learning careers saying, yeah, I'm donating 50 percent and my life is still like, absolutely awesome, in fact, is better as a result of the amount and donating. That's the sweet spot I think that you want to hit.

[02:09:57.720]
There's another issue here around how public to be around one's giving. And so, you know, you and I are having a public conversation about all of this. And this is just, by its very nature, violating a a norm that we've all inherited or a norm or a pseudo norm around generosity and altruism, which suggests that the highest form of generosity is to give anonymously. There's a Bible verse around this one where your virtue on your sleeve, you don't want to advertise your generosity because that conveys this message that you're doing it for reasons of self aggrandizement.

[02:10:41.730]
You're doing it to enhance your reputation. You want your name on the side of the building, whereas if you were really just connected to the cause of doing good, you would do all of this silently and people would find out after your death, or maybe they would never find out that you were the one who had secretly donated millions of dollars to cure some terrible disease or to buy bed nets. And yet, you know, you and I buy by association here, have flip that that ethic on its head, because it seems to be important to change people's thinking around all of the issues we've been discussing.

[02:11:22.230]
And the only way to do that is to really discuss them. And what's more, we're leveraging a concern about reputation kind of from the opposite side in recognizing that taking a pledge has psychological consequences. Right? I mean, when you publicly commit to do something that not only advertises to people that this is the sort of project a human being can become enamored of, you then have a reputational cost to worry about if you decide that you're going to renege on your your offer.

[02:11:53.370]
So talk for a few minutes about the significance of talking about any of this in the first place. Yeah.

[02:12:01.410]
So I think the public aspect is very important and it's for the reasons you mentioned earlier, that take the amount of goods that you're going to do in your life via donations and then just think, can I convince one other person to do the same? If so, you've doubled your impact. You've done your life's work over again. And I think possibly people can do that many times over, at least in the world today by being this kind of inspirational role model for others.

[02:12:27.150]
And so I think this religious tradition where, no, you shouldn't show the generosity you're doing, you should keep that secret. I think that looks pretty bad from an outcome, an outcome oriented perspective. And I think you need to be careful about how you're doing it. You want to be effective in your communication as well as your giving where, you know, it was very notable that Peter Singer had these arguments and giving for almost four decades with comparatively little uptake, certainly compared to the last 10 years of the effective altruism movement.

[02:12:59.190]
And, you know, my best hypothesis is that move them a framing that appeals primarily to guilt, which

is, you know, it's a low that I was motivation. You don't often get up and start really doing things on the basis of guilt to inspiration, instead saying like, no, this is amazing opportunity we have. And so this is a norm that I just really want to change. You know, in the long run, I would like it to be a part of common sense morality that use a significant part of your resources to help other people.

[02:13:29.010]
And we will only get that we'll only have that sort of cultural change if people are public about what they're doing and able to say, yeah, this is something I'm doing. I'm proud of it. I think you should consider doing it, too. This is the world I want to see. Well, well, you have certainly gotten the ball rolling in my life, and it's something that I'm immensely grateful for and I think this is a good place to leave it.

[02:13:52.860]
I know there will be questions and perhaps we can build out further lessons just based on frequently asked questions that come in in response to what we've said here. But I think that'll be the right way to proceed. So for the meantime, thank you for doing this, because you I think you're you're aware of how many people you're affecting, but it's still early days. And, you know, I think it'll be very interesting to see where all this goes, because I know what it's like to experience a tipping point around these issues personally.

[02:14:27.240]
And I have to think that that many people listening to us will have a similar experience one day or another and you will have occasion to it. So thank you for what you're doing.

[02:14:37.290]
Well, thank you for taking the pledge and getting involved. And yeah, I'm excited to see how these ideas develop over the coming years.